

Lab #8.1: Create a database and set up tables

Database: Library, database implementation

Overview

Objective

Use basic SQL DDL statements to set up a database and tables in a database. Based on the logical database design for a business:

- Interpret a logical ERD.
- Finalize the physical design.
- Create the database for the library.
- Create the tables.
- Populate the tables.
- Verify the database setup.

Outcomes

- Setup a database with all tables.
- Populate all tables with records.

Knowledge expected

- Create a database using SQL syntax.
- Create tables using SQL syntax.
- Work with the following `psql` commands: `\q`, `\?`, `\c`, `\l`, `\dt`, `\d [table]`
- To execute SQL statements from a file with the `psql \i` sub-command.

Submission instructions

- **READ ALL THE WORDS**
- **You must follow ALL submission instructions:** Submission instructions are explained in the “Lab 8 submission details” document, posted on BrightSpace. ANY submission instructions below NOT followed result in a grade of zero.
Note: This includes extra functionality that has not been requested.
- You are expected to complete all exercises, even those not required to be submitted.

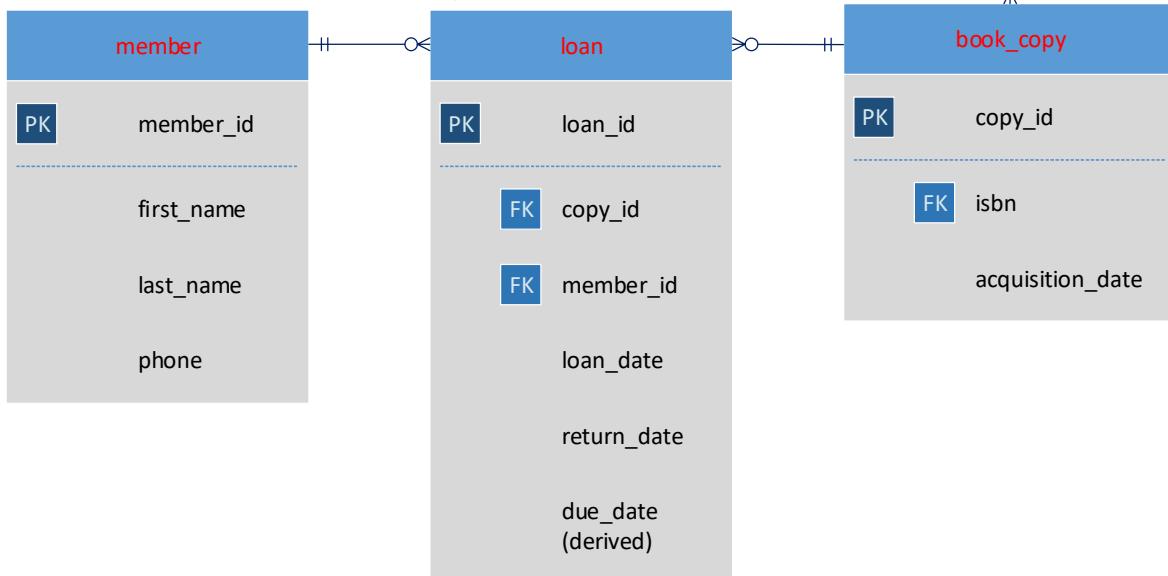
Background

Logical design for relational library database

Library ERD: Logical design

Note: If there is a discrepancy between the logical ERD and the lab description **the ERD is authoritative.**

Note on loan: A separate primary_key (loan_id) rather than a composite key of the foreign keys (copy_code + member_id) enables the database to store multiple loans of the same book by the same member.



Section A - Physical design

- The most functional data types for the attributes of the library relations are listed below:
 - int or smallint: integer (4-byte) or 2-byte integer
Example: rental_days: smallint
Note: The PostgreSQL data type serial is an integer that is auto-incremented by the database: it is used to generate surrogate primary keys.
 - char (length): where length is a fixed number of characters
Example: isbn: char(13)
 - varchar (maximum_length): variable number of characters up to maximum_length
Example: email: varchar(200)
 - date: date
Example: borrow_date: date

- Below find the completed table that implements the physical design by specifying the datatype for each attribute:

Entity	Relation name	Attribute names - identify PK - identify FK	Data type
Book	book	PK: isbn - title - rental_days	char(13) varchar(255) smallint
Book copy	book_copy	PK: copy_id FK: isbn - acquisition_date	int -> serial char(13) date
Member	member	PK: member_id - first_name - last_name - phone	int -> serial varchar(50) varchar(200) varchar(20)
Book loan	loan	PK: loan_id FK: copy_id FK: member_id - loan_date - return_date - due_date (derived)	int -> serial int int date date n/a: NOT stored in database

SQL overview

SQL syntax

- The semicolon at the end of the statement is part of the SQL syntax.
- SQL commands are case-insensitive.

Common syntax failures

The `psql` prompt gives an indication if there is a SQL syntax failure.

- of a regular prompt for a superuser: `dabase_name=#`
- Example* of a regular prompt for a non-superuser: `dabase_name=>`
- Example* of prompt indicating a missing closing bracket: `dabase_name (#`
- Example* of prompt indicating a missing closing semi-colon: `dabase_name-#`
- To cancel a bad command use: `\r`.

Section B – Create a database

Syntax practice #1: Create a new database using SQL syntax.

To create a database as `dbadmin` with the SQL `CREATE DATABASE` statement:

- Connect to the **postgres** database as the linux user **dbadmin**.
- At the **psql** prompt type: `CREATE DATABASE lib_your_network_ID;`
 - **Example:** `CREATE DATABASE lib_smit0001;`
 - **Note:** Don't forget the semicolon at the end of the statement – it is part of the syntax.
 - **Note:** To remove a database, use: `DROP DATABASE database`
- List all databases to verify that it has been created using the **psql \l** sub-command.

Switch to the newly created database using the subcommand `\c database`: the **psql** prompt displays the name of the database you are currently connected to.

Section C – Create tables

Overview of database creation options

- We will setup tables for the new database using SQL statements.
Note: To drop a table: `DROP TABLE table_name;`
- The database accepts SQL statements from the **psql** prompt or by importing a text file that contains SQL statements.
- Either method is acceptable: the instructions given here are for Linux text file creation and importing. To enter SQL statements on the **psql** prompt see the section labeled “Alternative method” at the end of the lab.

Steps to setup a table

- To create a table follow the physical design:
 - Use the proper naming convention for table and attribute names.
 - Use appropriate data types.
 - Set the primary key.
 - Set up foreign key references where required.
 - Decide when to use the NOT NULL (and optionally other) constraint.
- Create tables in proper order: a table has to be created before it can be referenced.
- Verify the successful creation of each table in the database.

Subcommands of **psql** to verify table creation

The interactive **psql** utility has a sub-command used to list all tables.

- To list all tables: `\dt`
Note: `\d` lists tables, views (virtual tables – *not covered*), and sequences (created for each serial data type).
- To list the fields of a table: `\d table_name`.

- Lists attributes, data types and modifiers
- Lists indexes
- Lists foreign key constraints
- **Example:** \d member

Preparation of data

- Create a library directory named `library_db` in the home directory of the user account responsible for managing the library database, dbadmin.
- In your library directory, create a subdirectory named `sql_create_table.d`.
- Within that directory set up one text file for each table and name it `create_<table_name>.sql`.
- **Example:** `~/library_db/sql_create_table.d/create_book.sql`

Verify that you are connected to the newly created database **before** creating tables.

Syntax practice #2: Create tables using SQL syntax.

Each file in the `sql_create_table.d` directory contains the SQL statement to create one table.

Table: member

Set the following constraints: PK, NOT NULL

```
CREATE TABLE member (
    member_id      serial PRIMARY KEY,
    first_name     varchar(100) NOT NULL,
    last_name      varchar(200) NOT NULL,
    phone          varchar(20)
);
```

Table: book

Set the following constraints: PK, NOT NULL

```
CREATE TABLE book (
    isbn           char(13) PRIMARY KEY,
    title          varchar(255) NOT NULL,
    rental_days    int NOT NULL
);
```

Table: book_copy

Set the following constraints: PK, FK referencing the book table.

```
CREATE TABLE book_copy (
    copy_id        serial PRIMARY KEY,
```

```

isbn          char(13) REFERENCES book(isbn),
acquisition_date date
);

```

Syntax practice #3: Execute SQL statements from a file and verify the result.

Import all tables into the library database and verify the result.

- The interactive `psql` utility has a sub-command used to execute SQL statements from a file.
 - **Syntax:** `\i file`
 - **Example:** `\i create_member.sql`
Note: This assumes that you started `psql` from the directory that contain the files; alternative use the Linux path statement to locate the file in the Linux directory structure.
- Verify the successful creation of all tables in the library database, including all relevant columns, data types and constraints.

Exercise #4: Create tables to complete the database design requirements.

Follow the steps below to create the table for `loan` based on the physical design.

- Use SQL syntax to create the `loan` table in the file in the `create_loan.sql` with all required attributes.
- In addition, set the following constraints:
 - Set the primary key.
 - Set a foreign key that references the `book_copy` table.
Note: When referencing an attribute of data type `serial`, associate the data type `integer` with the foreign key attribute.
 - Set a foreign key that references the `member` table.
- Import the file into the database: this executes the SQL statement.
- Verify the successful creation of all tables in the library database, including all relevant columns, data types and constraints.

Alternative method: SQL statement on the `psql` prompt

- To enter a multi-line SQL statement:
 - Statements in `psql` can span multiple lines: the ‘`;`’ is the “end-of-statement” marker.
 - To enter an SQL statement over multiple lines you can press [Enter] to split up the statement into multiple lines.
Note: You will notice that *the prompt changes* when entering a multi-line SQL statement.

- Once a table has been created successfully `psql` displays `CREATE TABLE`. Else, errors are identified explicitly.
- FYI:* After a table definition is entered, `psql` provides “notices” to inform you of additional actions it performs on the table.

Example for book_copy table:

- Creates a ‘sequence’ for the primary key because we are using the `serial` (auto-incrementing integer) data type.
- Creates an ‘index’ on the primary key.